

Case Study

In [1]:

```
# Manipulation
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns
```

Part 1: Data Exploration and Evaluation

1.1 Load and Explore

In [2]:

```
df = pd.read_csv('loan.csv', usecols=['loan_amnt',
                                       'funded_amnt',
                                       'term',
                                       'int_rate',
                                       'grade',
                                       'annual_inc',
                                       'issue_d',
                                       'dti',
                                       'revol_bal',
                                       'total_pymnt',
                                       'loan_status'])
```

In [3]:

```
df.shape # get shape
```

Out[3]:

```
(887379, 11)
```

In [4]:

```
df.head(10) # get first 10 rows
```

Out[4]:

| | loan_amnt | funded_amnt | term | int_rate | grade | annual_inc | issue_d | loan_status | d |
|---|-----------|-------------|-----------|----------|-------|------------|----------|-------------|------|
| 0 | 5000.0 | 5000.0 | 36 months | 10.65 | B | 24000.0 | Dec-2011 | Fully Paid | 27.6 |
| 1 | 2500.0 | 2500.0 | 60 months | 15.27 | C | 30000.0 | Dec-2011 | Charged Off | 1.0 |
| 2 | 2400.0 | 2400.0 | 36 months | 15.96 | C | 12252.0 | Dec-2011 | Fully Paid | 8.7 |
| 3 | 10000.0 | 10000.0 | 36 months | 13.49 | C | 49200.0 | Dec-2011 | Fully Paid | 20.0 |
| 4 | 3000.0 | 3000.0 | 60 months | 12.69 | B | 80000.0 | Dec-2011 | Current | 17.9 |
| 5 | 5000.0 | 5000.0 | 36 months | 7.90 | A | 36000.0 | Dec-2011 | Fully Paid | 11.2 |
| 6 | 7000.0 | 7000.0 | 60 months | 15.96 | C | 47004.0 | Dec-2011 | Current | 23.5 |
| 7 | 3000.0 | 3000.0 | 36 months | 18.64 | E | 48000.0 | Dec-2011 | Fully Paid | 5.9 |
| 8 | 5600.0 | 5600.0 | 60 months | 21.28 | F | 40000.0 | Dec-2011 | Charged Off | 5.5 |
| 9 | 5375.0 | 5375.0 | 60 months | 12.69 | B | 15000.0 | Dec-2011 | Charged Off | 18.0 |

In [5]:

```
df.sample(10) # get 10 random rows
```

Out[5]:

| | loan_amnt | funded_amnt | term | int_rate | grade | annual_inc | issue_d | loan_status |
|--------|-----------|-------------|-----------|----------|-------|------------|----------|-----------------|
| 519837 | 12000.0 | 12000.0 | 36 months | 9.17 | B | 40000.0 | Nov-2015 | Current |
| 380439 | 21000.0 | 21000.0 | 36 months | 8.39 | A | 80000.0 | May-2014 | Current |
| 883813 | 15000.0 | 15000.0 | 36 months | 15.59 | D | 55000.0 | Jan-2015 | Current |
| 541289 | 18325.0 | 18325.0 | 36 months | 13.67 | C | 53173.2 | Nov-2015 | Current |
| 252175 | 7700.0 | 7700.0 | 36 months | 13.66 | C | 41600.0 | Nov-2014 | Current |
| 262792 | 30000.0 | 30000.0 | 60 months | 12.99 | C | 71300.0 | Nov-2014 | Current |
| 808013 | 28000.0 | 28000.0 | 36 months | 7.89 | A | 144000.0 | Mar-2015 | Current |
| 490597 | 10000.0 | 10000.0 | 60 months | 18.49 | E | 68000.0 | Dec-2015 | Current |
| 288999 | 31075.0 | 31075.0 | 60 months | 24.50 | F | 82000.0 | Oct-2014 | Current |
| 727183 | 20000.0 | 20000.0 | 60 months | 13.99 | C | 70000.0 | Jun-2015 | In Grace Period |

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 887379 entries, 0 to 887378
Data columns (total 11 columns):
loan_amnt      887379 non-null float64
funded_amnt    887379 non-null float64
term           887379 non-null object
int_rate       887379 non-null float64
grade          887379 non-null object
annual_inc     887375 non-null float64
issue_d        887379 non-null object
loan_status    887379 non-null object
dti            887379 non-null float64
revol_bal      887379 non-null float64
total_pymnt    887379 non-null float64
dtypes: float64(7), object(4)
memory usage: 74.5+ MB
```

- There is a date variable: I will convert it to the correct type

In [7]:

```
df['issue_d'] = pd.to_datetime(df['issue_d'])
```

In [8]:

```
# Check  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 887379 entries, 0 to 887378  
Data columns (total 11 columns):  
loan_amnt      887379 non-null float64  
funded_amnt    887379 non-null float64  
term           887379 non-null object  
int_rate       887379 non-null float64  
grade          887379 non-null object  
annual_inc     887375 non-null float64  
issue_d        887379 non-null datetime64[ns]  
loan_status    887379 non-null object  
dti            887379 non-null float64  
revol_bal      887379 non-null float64  
total_pymnt    887379 non-null float64  
dtypes: datetime64[ns](1), float64(7), object(3)  
memory usage: 74.5+ MB
```

1.2 Look at different levels of aggregations

In [9]:

```
means = df.groupby(['term', 'grade', 'loan_status', df['issue_d'].dt.year]).mean()
```

In [10]:

```
means.head()
```

Out[10]:

| | | | | loan_amnt | funded_amnt | int_rate | annual_inc | | |
|-----------|-------|-------------|---------|-------------|-------------|----------|--------------|---|--|
| term | grade | loan_status | issue_d | | | | | | |
| 36 months | A | Charged Off | 2007 | 2000.000000 | 2000.000000 | 7.750000 | 24000.000000 | | |
| | | | 2008 | 8211.764706 | 7357.352941 | 8.345294 | 53167.764706 | | |
| | | | 2009 | 7457.594937 | 7457.594937 | 8.916203 | 55953.139241 | 1 | |
| | | | 2010 | 8045.045045 | 7545.945946 | 7.205856 | 54519.585586 | 1 | |
| | | | 2011 | 7599.719888 | 7584.453782 | 7.382689 | 51672.414482 | 1 | |

In [11]:

```
means.tail()
```

Out[11]:

| | | | | loan_amnt | funded_amnt | int_rate | annual_in |
|-----------|-------|--------------------|---------|--------------|--------------|-----------|--------------|
| term | grade | loan_status | issue_d | | | | |
| 60 months | G | Late (31-120 days) | 2011 | 14600.000000 | 14600.000000 | 21.590000 | 185000.00000 |
| | | | 2012 | 23000.000000 | 23000.000000 | 24.795000 | 67500.00000 |
| | | | 2013 | 24542.592593 | 24542.592593 | 25.557407 | 97346.22222 |
| | | | 2014 | 21221.785714 | 21221.785714 | 25.882000 | 77361.82000 |
| | | | 2015 | 21901.495726 | 21901.495726 | 26.670940 | 74822.00136 |

1.3 Check for Missing values and or repeated rows

In [12]:

```
df[df.duplicated()==True]
```

Out[12]:

| loan_amnt | funded_amnt | term | int_rate | grade | annual_inc | issue_d | loan_status | dti | rev |
|-----------|-------------|------|----------|-------|------------|---------|-------------|-----|-----|
|-----------|-------------|------|----------|-------|------------|---------|-------------|-----|-----|

Good news: There are no duplicated rows!

In [13]:

```
df.isnull().sum()
```

Out[13]:

```
loan_amnt      0
funded_amnt    0
term           0
int_rate       0
grade          0
annual_inc      4
issue_d        0
loan_status    0
dti            0
revol_bal      0
total_pymnt    0
dtype: int64
```

Only one column, annual income has missing values that too only 4 values.

- Now let's explore those missing values:

In [14]:

```
df[df.isnull().any(axis=1)]
```

Out[14]:

| | loan_amnt | funded_amnt | term | int_rate | grade | annual_inc | issue_d | loan_status |
|-------|-----------|-------------|-----------|----------|-------|------------|------------|---|
| 42449 | 5000.0 | 5000.0 | 36 months | 7.43 | A | NaN | 2007-08-01 | Does not meet the credit policy. Status:Fully ... |
| 42450 | 7000.0 | 7000.0 | 36 months | 7.75 | A | NaN | 2007-08-01 | Does not meet the credit policy. Status:Fully ... |
| 42480 | 6700.0 | 6700.0 | 36 months | 7.75 | A | NaN | 2007-07-01 | Does not meet the credit policy. Status:Fully ... |
| 42533 | 6500.0 | 6500.0 | 36 months | 8.38 | A | NaN | 2007-06-01 | Does not meet the credit policy. Status:Fully ... |

- Compare with any random 4 rows to see the difference

In [15]:

```
df.sample(4) # get 4 random rows
```

Out[15]:

| | loan_amnt | funded_amnt | term | int_rate | grade | annual_inc | issue_d | loan_status |
|--------|-----------|-------------|-----------|----------|-------|------------|------------|-------------|
| 423496 | 17000.0 | 17000.0 | 36 months | 13.65 | C | 90000.0 | 2014-03-01 | Current |
| 704441 | 2200.0 | 2200.0 | 36 months | 12.29 | C | 63000.0 | 2015-07-01 | Current |
| 431981 | 7000.0 | 7000.0 | 36 months | 10.99 | B | 99000.0 | 2014-03-01 | Fully Paid |
| 536790 | 14000.0 | 14000.0 | 36 months | 7.26 | A | 57000.0 | 2015-11-01 | Current |

- Loan status looks *interesting* here for the **rows with missing annual income...**

In [16]:

```
df[df.isnull().any(axis=1)][ 'loan_status' ]
```

Out[16]:

```
42449    Does not meet the credit policy. Status:Fully ...
42450    Does not meet the credit policy. Status:Fully ...
42480    Does not meet the credit policy. Status:Fully ...
42533    Does not meet the credit policy. Status:Fully ...
Name: loan_status, dtype: object
```

In [17]:

```
df[ 'loan_status' ].unique()
```

Out[17]:

```
array(['Fully Paid', 'Charged Off', 'Current', 'Default',
      'Late (31-120 days)', 'In Grace Period', 'Late (16-30 days)',
      'Does not meet the credit policy. Status:Fully Paid',
      'Does not meet the credit policy. Status:Charged Off', 'Issued'],
      dtype=object)
```

My evaluation

- Since these are less than **0.001%** of the rows, we can just drop them
- Imputing them doesn't make sense here since that will leak data

In [18]:

```
df_no_na = df.dropna(axis = 0)
```

In [19]:

```
df_no_na.shape # we have dropped those 4 rows
```

Out[19]:

```
(887375, 11)
```

1.3 Describe the distribution of features

1.3.1 Continuous features

In [20]:

```
df_cont = df_no_na.select_dtypes(include = [ "float64" ])
```

In [21]:

```
df_cont.sample(5)
```

Out[21]:

| | loan_amnt | funded_amnt | int_rate | annual_inc | dti | revol_bal | total_pymnt |
|--------|-----------|-------------|----------|------------|-------|-----------|-------------|
| 49483 | 5000.0 | 5000.0 | 6.62 | 120000.0 | 9.79 | 11359.0 | 5222.92 |
| 838711 | 23925.0 | 23925.0 | 21.67 | 75000.0 | 36.53 | 48272.0 | 6591.90 |
| 702975 | 3500.0 | 3500.0 | 11.53 | 125000.0 | 29.49 | 7782.0 | 3698.15 |
| 333877 | 10000.0 | 10000.0 | 8.39 | 80000.0 | 12.69 | 5337.0 | 5357.89 |
| 471437 | 35000.0 | 35000.0 | 9.17 | 93000.0 | 29.06 | 24962.0 | 0.00 |

In [22]:

```
# Let's look at the summary statistics first
df_cont.describe()
```

Out[22]:

| | loan_amnt | funded_amnt | int_rate | annual_inc | dti | rev |
|-------|---------------|---------------|---------------|--------------|---------------|---------|
| count | 887375.000000 | 887375.000000 | 887375.000000 | 8.873750e+05 | 887375.000000 | 8.87375 |
| mean | 14755.302719 | 14741.915678 | 13.246764 | 7.502759e+04 | 18.157113 | 1.69208 |
| std | 8435.455353 | 8429.897443 | 4.381862 | 6.469830e+04 | 17.190629 | 2.24268 |
| min | 500.000000 | 500.000000 | 5.320000 | 0.000000e+00 | 0.000000 | 0.00000 |
| 25% | 8000.000000 | 8000.000000 | 9.990000 | 4.500000e+04 | 11.910000 | 6.44350 |
| 50% | 13000.000000 | 13000.000000 | 12.990000 | 6.500000e+04 | 17.650000 | 1.18750 |
| 75% | 20000.000000 | 20000.000000 | 16.200000 | 9.000000e+04 | 23.950000 | 2.08290 |
| max | 35000.000000 | 35000.000000 | 28.990000 | 9.500000e+06 | 9999.000000 | 2.90483 |

In [23]:

```
# Plot distributions of continuous features
```

```
fig, ax = plt.subplots(nrows = 3, ncols = 3, figsize = (20, 20))
```

```
for i, ax in enumerate(ax.ravel()):
```

```
    if i > 6:
```

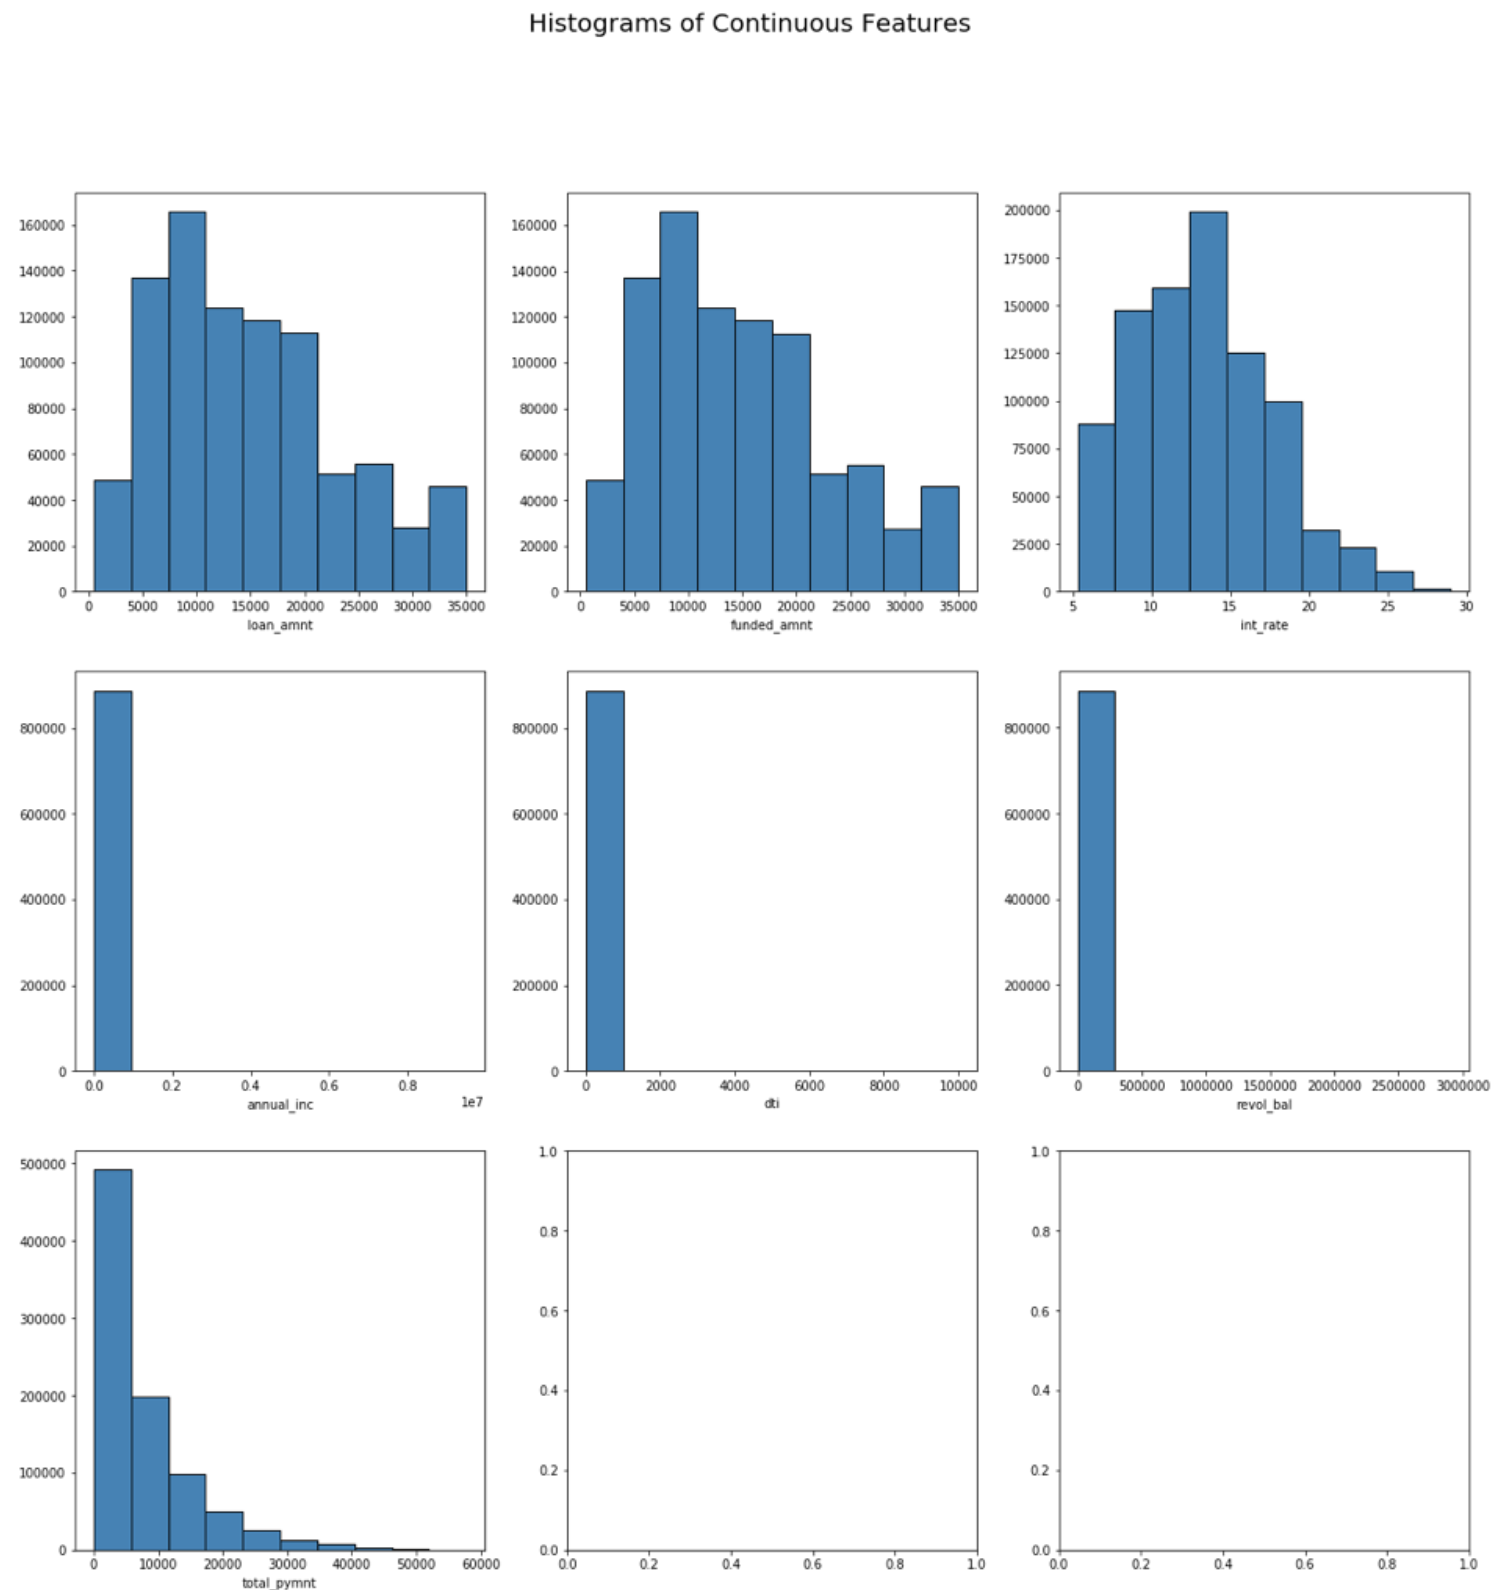
```
        continue
```

```
    ax.hist(df_cont.iloc[:, i], color='steelblue', edgecolor='black', bins=10,  
linewidth=1.0)
```

```
    ax.set_xlabel(df_cont.columns.values[i], color='k')
```

```
fig.suptitle('Histograms of Continuous Features', fontsize=20)
```

```
plt.show()
```



Comment: The distributions of the features are defined below

- Loan amount, Funded amount and Interest rate seem to have a slight positive skew; however, we can approximate them to be normally distributed
- Total payment: Has an exponential distribution, higher values have exponentially lower counts
- Annual income, DTI and Revolv_bal also have exponential distributions, with very high values of lambda, i.e. the decay is much faster

1.3.2 Categorical features

- Let's look at the unique values of the categorical columns

In [24]:

```
# separate categorical features
df_cat = df_no_na.select_dtypes(include = ["object"])
```

In [25]:

```
df_cat.sample(5)
```

Out[25]:

| | term | grade | loan_status |
|--------|-----------|-------|-------------|
| 687465 | 36 months | D | Current |
| 699920 | 60 months | G | Current |
| 770543 | 36 months | B | Current |
| 830349 | 60 months | B | Current |
| 211313 | 36 months | A | Fully Paid |

In [26]:

```
df_cat['term'].unique()
```

Out[26]:

```
array([' 36 months', ' 60 months'], dtype=object)
```

In [27]:

```
df_cat['grade'].unique()
```

Out[27]:

```
array(['B', 'C', 'A', 'E', 'F', 'D', 'G'], dtype=object)
```

In [28]:

```
df_cat['loan_status'].unique()
```

Out[28]:

```
array(['Fully Paid', 'Charged Off', 'Current', 'Default',  
      'Late (31-120 days)', 'In Grace Period', 'Late (16-30 days)',  
      'Does not meet the credit policy. Status:Fully Paid',  
      'Does not meet the credit policy. Status:Charged Off', 'Issued'],  
      dtype=object)
```

- **Summary statistics**

In [29]:

```
# Value Counts  
for col in df_cat:  
    print("Column Name: ", col )  
    print(" ")  
    print(df_cat[col].value_counts())  
    print(" ")
```

Column Name: term

36 months 621121

60 months 266254

Name: term, dtype: int64

Column Name: grade

B 254535

C 245860

A 148198

D 139542

E 70705

F 23046

G 5489

Name: grade, dtype: int64

Column Name: loan_status

Current 601779

Fully Paid 207723

Charged Off 45248

Late (31-120 days) 11591

Issued 8460

In Grace Period 6253

Late (16-30 days) 2357

Does not meet the credit policy. Status:Fully Paid 1984

Default 1219

Does not meet the credit policy. Status:Charged Off 761

Name: loan_status, dtype: int64

- **Visualisations**

In [30]:

```
# Plot distributions of continuous features

fig, ax = plt.subplots(nrows = 1, ncols = 3, figsize = (15, 5))
sns.set(style="darkgrid")


sns.countplot(x="term", data = df_cat, ax=ax[0])
sns.countplot(x="grade", data = df_cat, ax=ax[1])

# Set tick labels for loan status separatly since they don't fit
labels = df_cat['loan_status'].unique()
ax[2].set_xticklabels(labels = labels, rotation=90)

sns.countplot(x="loan_status", data = df_cat, ax=ax[2])

fig.suptitle('Count Plots of categorical features', fontsize=20)
plt.show()
```

Count Plots of categorical features



Comment

- We see that most of the loans are 36 month loans
- Grade B and C are the most common grades among the loans
- Similarly, for *loan status*, current is the most common status while the next common category is Fully paid followed by charged off. The rest of the categories are quite insignificant compared to these
- Check out the date variable

In [31]:

```
print("Min Date: ", df_no_na['issue_d'].min())
print("Min Date: ", df_no_na['issue_d'].max())
```

Min Date: 2007-06-01 00:00:00

Min Date: 2015-12-01 00:00:00

Outliers in the dataset

Given the skews in most of our continuous features, there exist sets of outliers; however, there is no point in removing them since I will be using an L2 penalty which reduces their effect on the model.

It is possible to understand outlier behaviour in this dataset using many approaches such as:

- IsolationForest
- OneClassSVM
- Elliptic Envelope

All of these are offered in sklearn and can be used to study these outliers; however, this is not the focus of this case study so will not be covered.

Part 2: Business Analysis

We are interested in evaluating whether the 36 month term loans would make for a good investment. Please investigate the following. Assume a 36 month investment period for each loan, and exclude loans with less than 36 months of data available.

In [32]:

```
# Subet dataframe
df_36_only = df_no_na[df_no_na['term'] == ' 36 months']
```

1) What percentage of loans has been fully paid?

In [33]:

```
perc_fully_paid = len(df_36_only[df_36_only['loan_status'] == 'Fully Paid']) /
len(df_36_only)
print("Percentage of loans Fully Paid: {:.2f}%".format(perc_fully_paid*100))
```

Percentage of loans Fully Paid: 26.98%

2) When bucketed by year of origination and grade, which cohort has the highest rate of defaults? Here you may assume that any loan which was not fully paid had “defaulted”.

- If 'Defaulted' = 1 i.e. anywhere loan not fully paid
- If 'Fully Paid' = 0

In [34]:

```
df_36_only['loan_status'] = np.where(df_36_only['loan_status'] == 'Fully Paid', 0, 1)
```

/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
 """Entry point for launching an IPython kernel.

In [35]:

```
# Check if the above code caused no errors
perc_fully_paid = len(df_36_only[df_36_only['loan_status'] == 0]) / len(df_36_only)
print("Percentage of loans Fully Paid {:.2f}%".format(perc_fully_paid*100))
```

Percentage of loans Fully Paid 26.98%

In [36]:

```
# Group by YEAR and grade
grouped = df_36_only.groupby([df_36_only['issue_d'].dt.year, 'grade'])
```

Rate of Default = Default / Total (In group)

In [37]:

```
# Apply function to get Default Rate

cohorts = grouped['loan_status'].agg([( 'perc_default',
                                         lambda x: ((x == 1).sum() / ((x == 0).sum() + (x == 1).sum())))]])
```

As is clear from the Dataframe below, 2007 G is the cohort with the highest default rate of basically 100%

In [38]:

```
# Sort values
cohorts.sort_values(by = ['perc_default'], ascending = False).head()
```

Out[38]:

| perc_default | | |
|--------------|-------|----------|
| issue_d | grade | |
| 2007 | G | 1.000000 |
| 2008 | G | 0.976744 |
| 2015 | A | 0.955312 |
| | B | 0.948912 |
| | C | 0.931936 |

3) When bucketed by year of origination and grade, what annualized rate of return have these loans generated on average?

- For simplicity, use the following approximation:
 - Annualized rate of return = $(\text{total_pymnt} / \text{funded_amnt})^{1/3} - 1$

In [39]:

```
averages = grouped.agg({'funded_amnt': 'mean', 'total_pymnt': 'mean'})
```

In [40]:

```
# define annualized rate of return
averages['Annualized_Rate_of_Returns'] = ((averages['total_pymnt'] / averages['funded_amnt'])**(1/3))-1
```

In [41]:

```
averages.head()
```

Out[41]:

| | | funded_amnt | total_pymnt | Annualized_Rate_of_Returns |
|---------|-------|-------------|-------------|----------------------------|
| issue_d | grade | | | |
| 2007 | A | 4945.945946 | 5396.874459 | 0.029511 |
| | B | 8184.693878 | 8323.797551 | 0.005633 |
| | C | 8132.978723 | 8245.527730 | 0.004592 |
| | D | 7654.040404 | 7540.993273 | -0.004948 |
| | E | 7817.750000 | 7633.371200 | -0.007924 |

Part 3: Modeling

Assumptions:

1. You are given the ability to invest in each loan independently; **i.e. IID Assumption**
2. You invest immediately following loan origination and hold to maturity (36 months); and
3. All loan fields that would be known upon origination are made available to you. **i.e. No missing data**

Discussion

- Class 1 = Default
- Class 0 = No Default

Task: Predict Loan Default

- As a risk-averse investor, I would like to ensure that those loans that I predict to be safe (i.e. will not default) **will actually not default** since I would like to ensure that my investment is safe
- Hence, we cannot afford **False Negatives**
- However, as an investor, I would like to be very sure of those I predict to default since I would not like to loose out on a perfectly good opportunity. Hence, **False Positives** are also costly.
- Hence, we would like to ensure that both false positives and false negatives are minimized
- The choice of metric will depend on the imbalance of the classes, which we know from part 2 is already imbalanced towards the "will default" class.
 - In this case metrics that are robust to class imbalance such as:
 - AUC
 - F1 Score: $2 \text{ (precision recall)} / (\text{precision} + \text{recall})$

Precision = True / Predicted to be True

- $tp / (tp + fp)$

Recall = True / Actual Truth

- $tp / (tp + fn)$

Data preparation for modelling

In [42]:

```
# Libraries for preprocessing
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder

# Modelling
from sklearn.linear_model import LogisticRegression

# Splitting and Cross-validation
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.model_selection import StratifiedShuffleSplit
from sklearn.model_selection import StratifiedKFold

# Metrics
from sklearn.metrics import accuracy_score, roc_auc_score, f1_score
from sklearn.metrics import confusion_matrix, classification_report
```

1. Drop missing data: Already dropped in the first part of this study

2. Separate Features and Target

In [43]:

```
y = df_36_only['loan_status']
X = df_36_only.drop(['loan_status'], axis = 1)
```

In [44]:

```
print("Default% = {:.2f}%".format((len(y[y==1]) / len(y))*100))
```

Default% = 73.02%

3. Drop columns that clearly don't add any value

In [45]:

```
X.shape
```

Out[45]:

```
(621121, 10)
```

In [46]:

```
X.sample(5)
```

Out[46]:

| | loan_amnt | funded_amnt | term | int_rate | grade | annual_inc | issue_d | dti | revol_bal |
|--------|-----------|-------------|-----------|----------|-------|------------|------------|-------|-----------|
| 710884 | 6000.0 | 6000.0 | 36 months | 12.69 | C | 99650.0 | 2015-06-01 | 30.96 | 7091.0 |
| 703362 | 16000.0 | 16000.0 | 36 months | 7.89 | A | 37500.0 | 2015-07-01 | 38.50 | 16000.0 |
| 705666 | 10000.0 | 10000.0 | 36 months | 12.69 | C | 43000.0 | 2015-07-01 | 13.90 | 11000.0 |
| 84676 | 8000.0 | 8000.0 | 36 months | 19.20 | D | 50000.0 | 2013-10-01 | 8.45 | 7000.0 |
| 408645 | 16000.0 | 16000.0 | 36 months | 12.49 | B | 104000.0 | 2014-04-01 | 13.44 | 19000.0 |

In [47]:

```
X.drop(['term', 'issue_d'], axis = 1, inplace =True)
```

In [48]:

```
X.sample(5)
```

Out[48]:

| | loan_amnt | funded_amnt | int_rate | grade | annual_inc | dti | revol_bal | total_pymnt |
|--------|-----------|-------------|----------|-------|------------|-------|-----------|-------------|
| 64084 | 6000.0 | 6000.0 | 8.90 | A | 60000.0 | 6.64 | 8691.0 | 6763.51000 |
| 226286 | 30000.0 | 30000.0 | 7.62 | A | 90000.0 | 0.66 | 4536.0 | 33201.88368 |
| 477501 | 24000.0 | 24000.0 | 6.89 | A | 120000.0 | 22.67 | 38405.0 | 0.00000 |
| 397187 | 18000.0 | 18000.0 | 8.90 | A | 64000.0 | 16.18 | 20886.0 | 19391.47000 |
| 312054 | 24925.0 | 24925.0 | 26.06 | G | 60000.0 | 33.72 | 56586.0 | 31888.33000 |

In [49]:

```
X.shape
```

Out[49]:

(621121, 8)

4. OneHotEncoding Categorical Variables

In [50]:

```
X_one_hot = pd.get_dummies(X)
```

5. Test Train Split: IID assumption

In [51]:

```
X_train, X_test, y_train, y_test = train_test_split(X_one_hot, y, test_size = 0.33)
```

6. Fit Logistic Regression Mode

- *penalty = l2* is default, it makes logistic less prone to overfitting (It's basically the ridge penalty)
- *C = 1* is the default hyperparameter value which is used to tune Logistic Regression
- *class_weight = balanced* mode uses the values of *y* to automatically adjust weights inversely proportional to class frequencies in the input data as $n_{\text{samples}} / (n_{\text{classes}} * \text{np.bincount}(y))$. **(We use this is because we have an imbalanced dataset)**
- *max_iter* used with solver == 'sag' (we need to vary the number of iterations to allow the to function to converge)
- *solver = 'sag'* corresponds to Stochastic Average Gradient and makes computation faster for large datasets such as this one
- *random_state* used with solver == 'sag' (since we're doing stochastic optimization)

Source: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html
(https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)

In [52]:

```
lr = LogisticRegression(penalty = 'l2',  
                        C = 1,  
                        class_weight = 'balanced',  
                        solver = 'sag',  
                        max_iter = 100,  
                        random_state = 4896)  
lr.fit(X_train, y_train)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/  
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac  
hed which means the coef_ did not converge  
  "the coef_ did not converge", ConvergenceWarning)
```

Out[52]:

```
LogisticRegression(C=1, class_weight='balanced', dual=False,  
                  fit_intercept=True, intercept_scaling=1, max_iter=100,  
                  multi_class='warn', n_jobs=None, penalty='l2', random_st  
ate=4896,  
                  solver='sag', tol=0.0001, verbose=0, warm_start=False)
```

In [53]:

```
y_pred = lr.predict(X_train)
```

In [54]:

```
# Evaluate model performance over these three metrics
print("Score on Train set (only)")
print("Accuracy score: {:.2f}".format(accuracy_score(y_train, y_pred)))
print("F1 score: {:.2f}".format(f1_score(y_train, y_pred)))
print("ROC AUC score: {:.2f}".format(roc_auc_score(y_train, y_pred)))
```

```
Score on Train set (only)
Accuracy score: 0.93
F1 score: 0.95
ROC AUC score: 0.95
```

In [55]:

```
cm = confusion_matrix(y_train, y_pred)
print("True Negatives:", cm[0][0])
print("False Negatives:", cm[1][0])
print("True Positives:", cm[1][1])
print("False Positives:", cm[0][1])
```

```
True Negatives: 111952
False Negatives: 28934
True Positives: 274892
False Positives: 373
```

- Our False Positive rate seems fine.
 - However, we seem to have quite a bit of False negatives in our case, which is **NOT GOOD**
-
- **Try scaling the numerical data to see if scores improve**

In [56]:

```
# Save continuous columns separately
X_train_cont = X_train.select_dtypes(include = ["float64"]).reset_index(drop =
True)
X_test_cont = X_test.select_dtypes(include = ["float64"]).reset_index(drop = T
rue)
```

In [57]:

```
# Save continuous column names
train_cont = list(X_train_cont)
test_cont = list(X_test_cont)
```

In [58]:

```
# Fit on X_train
scale = StandardScaler()
scale.fit(X_train_cont)
```

Out[58]:

```
StandardScaler(copy=True, with_mean=True, with_std=True)
```

In [59]:

```
# Save categorical columns separately
X_train_cat = X_train.select_dtypes(exclude = ["float64"]).reset_index(drop =
True)
X_test_cat = X_test.select_dtypes(exclude = ["float64"]).reset_index(drop = Tr
ue)
```

In [60]:

```
# Transform using X_train
X_train_scaled = scale.transform(X_train_cont)

# Transform using X_train to avoid information leakage
X_test_scaled = scale.transform(X_test_cont)
```

In [61]:

```
# Convert to Pandas Dataframe
X_train_scaled = pd.DataFrame(X_train_scaled, columns = train_cont).reset_inde
x(drop = True)
X_test_scaled = pd.DataFrame(X_test_scaled, columns = test_cont).reset_index(d
rop = True)
```

In [62]:

```
# Check
X_train_scaled.head()
```

Out[62]:

| | loan_amnt | funded_amnt | int_rate | annual_inc | dti | revol_bal | total_pymnt |
|---|-----------|-------------|-----------|------------|-----------|-----------|-------------|
| 0 | 1.610232 | 1.611412 | 0.301029 | -0.194231 | -0.079529 | -0.146563 | -0.977086 |
| 1 | -0.318240 | -0.317454 | 1.616361 | -0.318910 | 0.156310 | -0.167930 | 0.341735 |
| 2 | -1.089629 | -1.089000 | 0.070133 | -0.537097 | 0.254576 | -0.373046 | -0.922579 |
| 3 | -0.832499 | -0.831818 | -0.267132 | 0.012097 | 0.533993 | 0.776527 | -0.708784 |
| 4 | 1.555591 | 1.556761 | -0.809350 | -0.272155 | -0.427732 | 0.142902 | 2.124878 |

In [63]:

```
# Check
print(X_test_scaled.shape)
# Check
print(X_train_scaled.shape)
```

```
(204970, 7)
(416151, 7)
```

In [64]:

```
X_train_scaled = pd.concat([X_train_scaled, X_train_cat], axis = 1, join = "inner").reset_index(drop = True)
```

In [65]:

```
X_test_scaled = pd.concat([X_test_scaled, X_test_cat], axis = 1, join = "inner").reset_index(drop = True)
```

In [66]:

```
# Check
print(X_test_scaled.shape)
# Check
print(X_train_scaled.shape)
```

```
(204970, 14)
(416151, 14)
```

In [67]:

```
# Check
X_train_cat.head()
```

Out[67]:

| | grade_A | grade_B | grade_C | grade_D | grade_E | grade_F | grade_G |
|---|---------|---------|---------|---------|---------|---------|---------|
| 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 |
| 2 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| 3 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
| 4 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |

In [68]:

```
# Check
X_train_scaled.head()
```

Out[68]:

| | loan_amnt | funded_amnt | int_rate | annual_inc | dti | revol_bal | total_pymnt | grad |
|---|-----------|-------------|-----------|------------|-----------|-----------|-------------|------|
| 0 | 1.610232 | 1.611412 | 0.301029 | -0.194231 | -0.079529 | -0.146563 | -0.977086 | |
| 1 | -0.318240 | -0.317454 | 1.616361 | -0.318910 | 0.156310 | -0.167930 | 0.341735 | |
| 2 | -1.089629 | -1.089000 | 0.070133 | -0.537097 | 0.254576 | -0.373046 | -0.922579 | |
| 3 | -0.832499 | -0.831818 | -0.267132 | 0.012097 | 0.533993 | 0.776527 | -0.708784 | |
| 4 | 1.555591 | 1.556761 | -0.809350 | -0.272155 | -0.427732 | 0.142902 | 2.124878 | |

- **Fit model on scaled Training Data**

In [69]:

```
lr.fit(X_train_scaled, y_train)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
  "the coef_ did not converge", ConvergenceWarning)
```

Out[69]:

```
LogisticRegression(C=1, class_weight='balanced', dual=False,
                  fit_intercept=True, intercept_scaling=1, max_iter=100,
                  multi_class='warn', n_jobs=None, penalty='l2', random_st
ate=4896,
                  solver='sag', tol=0.0001, verbose=0, warm_start=False)
```

In [70]:

```
y_pred = lr.predict(X_train_scaled)
```

In [71]:

```
# Evaluate model performance over these three metrics
print("Score on Train set (only)")
print("Accuracy score: {:.2f}".format(accuracy_score(y_train, y_pred)))
print("F1 score: {:.2f}".format(f1_score(y_train, y_pred)))
print("ROC AUC score: {:.2f}".format(roc_auc_score(y_train, y_pred)))
```

```
Score on Train set (only)
Accuracy score: 0.94
F1 score: 0.96
ROC AUC score: 0.96
```


In [72]:

```
cm = confusion_matrix(y_train, y_pred)
print("True Negatives:", cm[0][0])
print("False Negatives:", cm[1][0])
print("True Positives:", cm[1][1])
print("False Positives:", cm[0][1])
```

True Negatives: 111903

False Negatives: 24096

True Positives: 279730

False Positives: 422


Scaling significantly improved performance

- Number of False Negatives fell by roughly 7000
- However, the Number for False Positives rose by 300

- **Look at feature importances**

In [73]:

```
n_features = X_train_scaled.shape[1]
coefs_lr = lr.coef_[0]
plt.barh(range(n_features), coefs_lr, align='center')
plt.yticks(np.arange(n_features), list(X_train_scaled.columns.values))
plt.xlabel("Feature importance")
plt.ylabel("Feature")
plt.ylim(-1, n_features)
plt.show()
```



7. Tune Parameters to optimize model performance: On Scaled Data

- Perform Grid-search on Sampled data to save time

In [74]:

```
# Cross-validation strategy: shuffles the data
cv = StratifiedShuffleSplit(n_splits = 5, random_state= 4896)
```

In [77]:

```
pipe_lr = make_pipeline(LogisticRegression(penalty = 'l2', class_weight = 'balanced', solver = 'sag', max_iter = 100,
                                             random_state = 4896))
param_grid_lr = [{'logisticregression__C':np.logspace(-3, 0, 5)}]
grid_lr = GridSearchCV(pipe_lr, param_grid_lr, scoring = 'roc_auc', cv = cv)
```

In [78]:

```
grid_lr.fit(X_train_scaled, y_train)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
```

[illegible]

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  "the coef_ did not converge", ConvergenceWarning)
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  "the coef_ did not converge", ConvergenceWarning)
```

Out[78]:

```
GridSearchCV(cv=StratifiedShuffleSplit(n_splits=5, random_state=4896, test_size='default',
    train_size=None),
    error_score='raise-deprecating',
    estimator=Pipeline(memory=None,
    steps=[('logisticregression', LogisticRegression(C=1.0, class_weight='balanced', dual=False,
    fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l2', random_state=4896,
    solver='sag', tol=0.0001, verbose=0, warm_start=False))]
),
    fit_params=None, iid='warn', n_jobs=None,
    param_grid=[{'logisticregression__C': array([0.001, 0.005, 0.01, 0.03162, 0.17783, 1.0])}],
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring='roc_auc', verbose=0)
```

In [79]:

```
print("Best Score %s" % grid_lr.best_score_)
print("Best Parameters %s" % grid_lr.best_params_)
```

```
Best Score 0.9811531963960071
Best Parameters {'logisticregression__C': 1.0}
```

In [80]:

```
best_param = grid_lr.best_params_.get('logisticregression__C')
```

Fit Best Parameter value

In [81]:

```
lr = LogisticRegression(penalty = 'l2',
                        C = best_param,
                        class_weight = 'balanced',
                        solver = 'sag',
                        max_iter = 100,
                        random_state = 4896)
lr.fit(X_train_scaled, y_train)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
  "the coef_ did not converge", ConvergenceWarning)
```

Out[81]:

```
LogisticRegression(C=1.0, class_weight='balanced', dual=False,
                  fit_intercept=True, intercept_scaling=1, max_iter=100,
                  multi_class='warn', n_jobs=None, penalty='l2', random_st
ate=4896,
                  solver='sag', tol=0.0001, verbose=0, warm_start=False)
```

In [82]:

```
y_pred = lr.predict(X_train_scaled)
```

In [83]:

```
# Evaluate model performance over these three metrics
print("Score on Train set (only)")
print("Accuracy score: {:.2f}".format(accuracy_score(y_train, y_pred)))
print("F1 score: {:.2f}".format(f1_score(y_train, y_pred)))
print("ROC AUC score: {:.2f}".format(roc_auc_score(y_train, y_pred)))
```

```
Score on Train set (only)
Accuracy score: 0.94
F1 score: 0.96
ROC AUC score: 0.96
```

In [84]:

```
cm = confusion_matrix(y_train, y_pred)
print("True Negatives:", cm[0][0])
print("False Negatives:", cm[1][0])
print("True Positives:", cm[1][1])
print("False Positives:", cm[0][1])
```

```
True Negatives: 111903
False Negatives: 24096
True Positives: 279730
False Positives: 422
```

8. Evaluating on Test Set and Check for overfitting

In [85]:

```
y_pred = lr.predict(X_test_scaled)
```

In [86]:

```
# Evaluate model performance over these three metrics
print("Score on Test set")
print("Accuracy score: {:.2f}".format(accuracy_score(y_test, y_pred)))
print("F1 score: {:.2f}".format(f1_score(y_test, y_pred)))
print("ROC AUC score: {:.2f}".format(roc_auc_score(y_test, y_pred)))
```

Score on Test set

Accuracy score: 0.94

F1 score: 0.96

ROC AUC score: 0.96

In [87]:

```
cm = confusion_matrix(y_test, y_pred)
print("True Negatives:", cm[0][0])
print("False Negatives:", cm[1][0])
print("True Positives:", cm[1][1])
print("False Positives:", cm[0][1])
```

True Negatives: 55057

False Negatives: 11900

True Positives: 137820

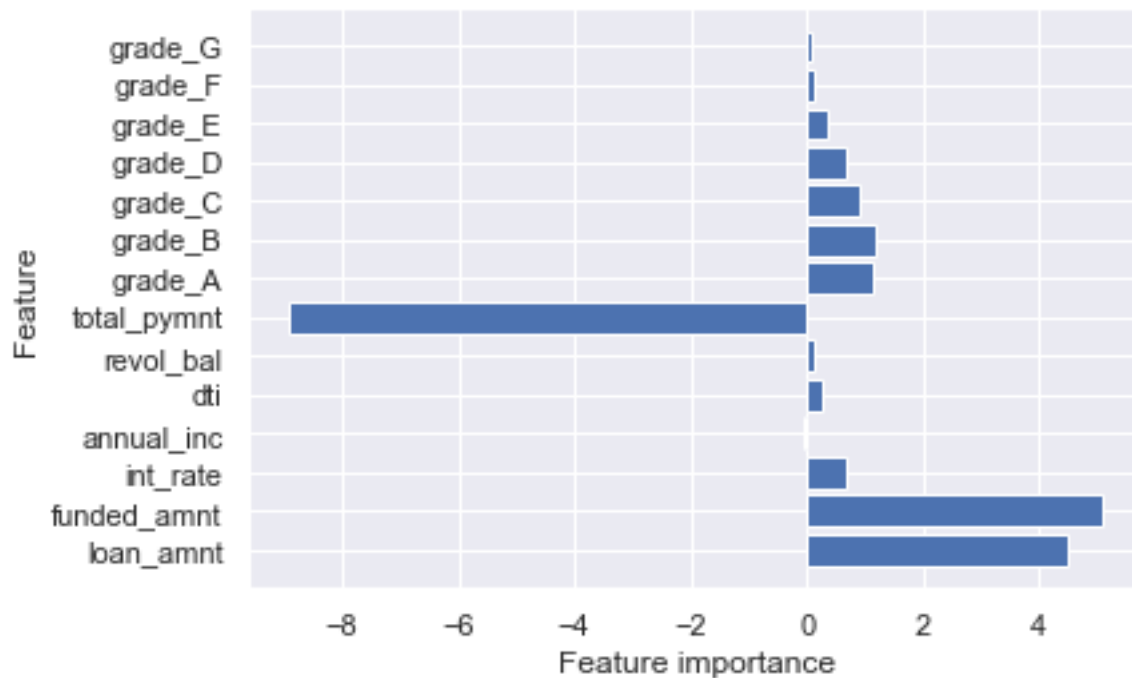
False Positives: 193

There seems to be no overfitting since train and test scores seem to be the similar.

9. Look at feature importances

In [88]:

```
n_features = X_train.shape[1]
coefs_lr = grid_lr.best_estimator_.named_steps["logisticregression"].coef_[0]
plt.barh(range(n_features), coefs_lr, align='center')
plt.yticks(np.arange(n_features), list(X_train.columns.values))
plt.xlabel("Feature importance")
plt.ylabel("Feature")
plt.ylim(-1, n_features)
plt.show()
```



10. Put everything in a Pipeline (To prevent any kind of data leakage)

In [89]:

```
# Scale Numeric Features

numeric_features = list(X.select_dtypes(include = ["float64"]))
numeric_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())])
```

In [90]:

```
# OneHotEncode Categorical Features

categorical_features = list(X.select_dtypes(exclude = ["float64"]))
categorical_transformer = Pipeline(steps=[
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])
```

In [91]:

```
# Add in preprocessor
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])
```


In [92]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.33)
```

In [93]:

```
# Define steps in Pipe
pipe = Pipeline(steps=[('preprocessor', preprocessor),
                        ('classifier', LogisticRegression(penalty = 'l2',
                                                         C = best_param,
                                                         class_weight = 'balanced',
                                                         solver = 'sag',
                                                         max_iter = 100,
                                                         random_state = 4896))]
) # solver='lbfgs'
```

Tune Parameters

In [98]:

```
param_grid = {
    'classifier__C': np.logspace(-3, 0, 5)
}

grid_search = GridSearchCV(pipe, param_grid, cv = cv)
grid_search.fit(X_train, y_train)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
```

```
"the coef_ did not converge", ConvergenceWarning)
```

```
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
```


[illegible]

```

/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
    "the coef_ did not converge", ConvergenceWarning)
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
    "the coef_ did not converge", ConvergenceWarning)
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
    "the coef_ did not converge", ConvergenceWarning)
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
    "the coef_ did not converge", ConvergenceWarning)
/Users/pranjalbajaj/anaconda3/lib/python3.7/site-packages/sklearn/
linear_model/sag.py:334: ConvergenceWarning: The max_iter was reac
hed which means the coef_ did not converge
    "the coef_ did not converge", ConvergenceWarning)

```

Out[98]:

```

GridSearchCV(cv=StratifiedShuffleSplit(n_splits=5, random_state=48
96, test_size='default',
        train_size=None),
        error_score='raise-deprecating',
        estimator=Pipeline(memory=None,
        steps=[('preprocessor', ColumnTransformer(n_jobs=None, remain
der='drop', sparse_threshold=0.3,
        transformer_weights=None,
        transformers=[('num', Pipeline(memory=None,
        steps=[('scaler', StandardScaler(copy=True, with_mean=True, w
ith_std=True))]), ['loan_amnt', 'funded_amnt', 'int... penalty='l2
', random_state=4896,
        solver='sag', tol=0.0001, verbose=0, warm_start=False))])
),
        fit_params=None, iid='warn', n_jobs=None,
        param_grid={'classifier__C': array([0.001 , 0.00562, 0.031
62, 0.17783, 1.      ])}},
        pre_dispatch='2*n_jobs', refit=True, return_train_score='wa
rn',
        scoring=None, verbose=0)

```

In [117]:

```
y_pred = pipe.predict(X_test)
```

In [118]:

```
# Evaluate model performance over these three metrics
print("Score on Test set")
print("Accuracy score: {:.2f}".format(accuracy_score(y_test, y_pred)))
print("F1 score: {:.2f}".format(f1_score(y_test, y_pred)))
print("ROC AUC score: {:.2f}".format(roc_auc_score(y_test, y_pred)))
```

Score on Test set

Accuracy score: 0.94

F1 score: 0.96

ROC AUC score: 0.96

In [119]:

```
cm = confusion_matrix(y_test, y_pred)
print("True Negatives:", cm[0][0])
print("False Negatives:", cm[1][0])
print("True Positives:", cm[1][1])
print("False Positives:", cm[0][1])
```

True Negatives: 54957

False Negatives: 11808


True Positives: 138004

False Positives: 201

Feature Importances

In [110]:

```
n_features = X_train_scaled.shape[1]
coefs_lr = grid_search.best_estimator_.named_steps['classifier'].coef_[0]
plt.barh(range(n_features), coefs_lr, align='center')
plt.yticks(np.arange(n_features), list(X_train_scaled))
plt.xlabel("Feature importance")
plt.ylabel("Feature")
plt.ylim(-1, n_features)
plt.show()
```



Was the model effective? Explain how you validated your model and describe how you measure the performance of the model.

Metrics

For this study, I chose two metrics to optimize performance:

- AUC
- F1 Score: $2 \times (\text{precision} \times \text{recall}) / (\text{precision} + \text{recall})$

Both of these metrics are fairly robust to class imbalance and capture our need to minimize:

- False Negatives
- False Positives

Since both these scenarios can be costly to an investor and should be minimized.

Model Performance

- The final scores on the test set were:

Metrics

- Score on Test set
- Accuracy score: 0.94
- F1 score: 0.96
- ROC AUC score: 0.96

Confusion Matrix Results

- True Negatives: 54957
- False Negatives: 11808
- True Positives: 138004
- False Positives: 201

After the pipeline, final scores on the test set were:

Overall, the model performs well according to these metrics; however, to truly gauge the performance of a model, it must be compared to a dummy classifier or a random baseline. See results below.

In [111]:

```
from sklearn.dummy import DummyClassifier
```

In [112]:

```
dummy = DummyClassifier(strategy='stratified', random_state=None, constant=None)
dummy.fit(X_train_scaled, y_train)
y_pred = dummy.predict(X_test_scaled)

# Evaluate model performance over these three metrics
print("Score on Test set")
print("Accuracy score: {:.2f}".format(accuracy_score(y_test, y_pred)))
print("F1 score: {:.2f}".format(f1_score(y_test, y_pred)))
print("ROC AUC score: {:.2f}".format(roc_auc_score(y_test, y_pred)))

print("Confusion Matrix Metrics")
cm = confusion_matrix(y_test, y_pred)
print("True Negatives:", cm[0][0])
print("False Negatives:", cm[1][0])
print("True Positives:", cm[1][1])
print("False Positives:", cm[0][1])
```

```
Score on Test set
Accuracy score: 0.61
F1 score: 0.73
ROC AUC score: 0.50
Confusion Matrix Metrics
True Negatives: 14815
False Negatives: 40346
True Positives: 109466
False Positives: 40343
```

In [120]:

```
print("F1 score improvement: {:.2f}%".format((0.96 - 0.73)*100 / 0.73))
print("AUC improvement: {:.2f}%".format((0.96 - 0.5)*100 / 0.5) )
```

```
F1 score improvement: 31.51%
AUC improvement: 92.00%
```

Compared to this dummy classifier, we seem to be doing much better in terms of both F1 Score and AUC score. I do have reason to believe that the model was effective.

Note on Feature Importances:

- It is key in an evaluation to draw upon feature importances and we find consistently that:
 - Total Payment, Funded Amount and Loan Amount in this order are the most important features
 - This makes complete sense since our EDA discussed revealed that these features were the most informative!

Note:

- I made sure to make predictions on the Test Set ONLY once i.e. after the model tuning stage.
- I created a new test and train split during the pipeline stage, which I used to generate another set of predictions on the test set once after all parameters were trained. This was done to ensure that the pipeline works.

Further comments


Given the context of this data, using an IID split will mean that we're using future values to present the past values which doesn't seem to make much sense in the real world. Further, effective EDA should involve looking at the temporal component of the data as well as below.

In [115]:

```
cohorts.reset_index(inplace = True)
```

In [116]:

```
plt.suptitle('Time Series of % Loans Defaulted', fontsize = 14)
plt.xlabel('Year', fontsize = 12)
plt.ylabel('% of Default', fontsize = 12)
plt.plot(cohorts['issue_d'], cohorts['perc_default'])
plt.show()
```



There clearly is a temporal trend in the data which makes it very difficult to assume that observations are IID.

- The recession period from 2007-9 clearly shows very high default deates, with rates slipping as they go down to 2012
- The reates begin to rise again after 2012 in a strange pattern

Next steps and improved modelling approaches:

- The approach should follow using, say 2007-2012 data for the train set and 2013-15 as the test set.
- To tune model parameters, it may be possible to use *TimeSeriesSplit* offered by sklearn
- Given that there were only 4 missing values, we wouldn't have gained much imputing them; however, if there were more missing values then using packages such as *fancyimpute* may make sense which impute continous data by imputing values that of similar rows
- Tree-based models such as Random Forest or XG Boost are often good performers, not prone to overfitting and equally interpretable as Logistic Regression thanks to **shap values** which allow for interpretibility of these complex models